

Task 2: Developing an Inheritance Hierarchy of Neural Networks

We utilize Tensorflow Keras API to establish a class hierarchy for image classification. At the base we have the 'tf.keras.Model'. Extending this we have the superclass 'NeuralNetwork' as a blueprint for neural networks, encapsulating common behaviors and properties. that again has two subclasses for different types of neural networks:

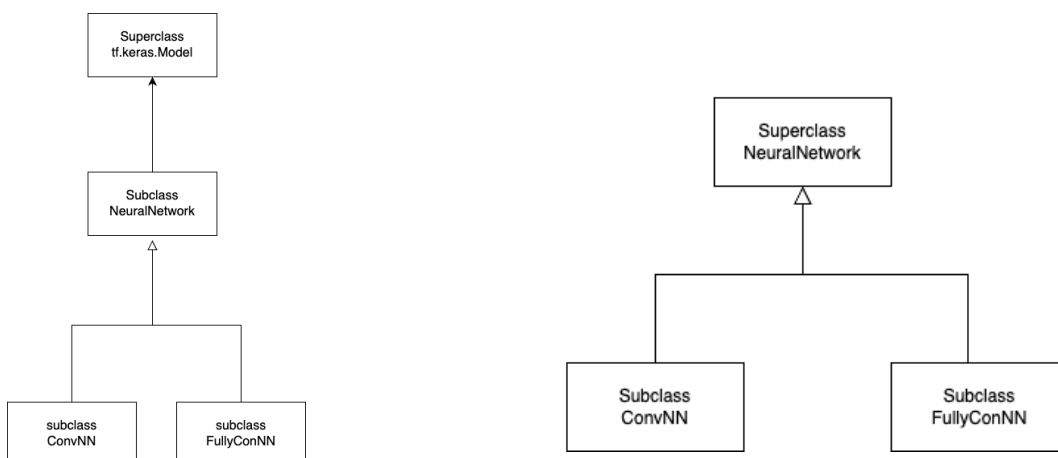
- Convolutional neural networks
- Fully connected neural networks

From 'NeuralNetwork' we derive the subclasses 'ConvNN' and 'FullyConNN'. The difference between these classes is that in 'ConvNN', hidden layers are set up to create a stack of convolutional layers and flatten it to convert the 2D feature maps into a 1D vector. In 'FullyConNN' the input is expected to be a flat vector (28*28) here, all the neurons in one layer are connected to the subsequent layer.

The Hierarchy of the classes has tf.keras.Model on top as the Superclass, this class is builtin and is called through the TensorFlow Keras API. This model handels training, prediction, saving, loading and more. NeuralNetwork is a subclass of tf.keras.Model, where you can define common behaviors and properties that are shared across various types of NNs. NeuralNetwork inherits from Keras and therefore gets all the functionality of the keras model, as well as the functionality you add FullyConNN and ConvNN, which are specialized versions of NNs. They are subclasses of the NeuralNetwork class. ConvNN overrides specific methods to set up the convolutional layers explained earlier. Similarly for FullyConNN, this overrides methods to set up densely connected layers.

2.1) Inheritance Diagram and Public Interface:

The NeuralNetwork class contains constructor `__init__` together with the methods: `__repr__`, `_hidden_layers`, `classifier`, `_call`, `train`, `test`, `set_hidden`, `set_cls` and `set_params`. FullyConNN contains the constructor `__init__` and method `_hidden_layers`. ConvNN contains the same as FullyConv: `__init__` and `_hidden_layers`.



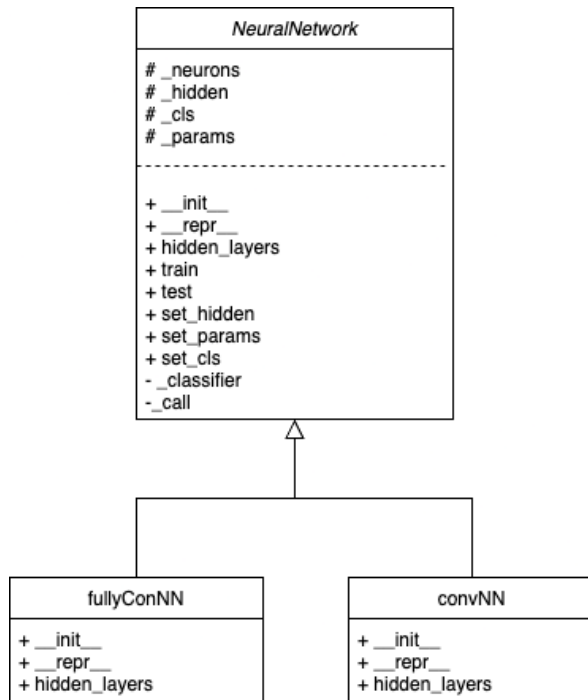
The illustration to the left shows us how NeuralNetwork inherits from the tf.keras.Model, and the illustration to the right show us how we imagine our class hierarchy, where the NeuralNetwork is the Superclass of ConvNN and FullyConNN

Public interface:

instance variables

+ Public methods

- Private methods



2.2) Implementation of code:

See python-files.

2.3) Inheritance, Overriding, and Polymorphism:

For full explanation of the parameters taken into methods, and the return of the methods, see docstring. In the code you can write `-doc <classname>` for docstring of a specific class.

The NeuralNetwork class:

Constructor: In the constructor method we chose to not take in any variables. This is because we assume that the values of the instance variables can be changed and set later.

We choose to have instance variables `_neurons`, `_hidden`, `_cls` and `_params`.

All set to None, and then we make them later when we call the `hidden_layers` method for the subclasses.

hidden_layers: Because of the fact that the `hidden_layers` method is implemented differently in each subclass we make it abstract in the superclass. Each subclass overrides the method.

This is to be able to use the same method call in our programme. The `hidden_layers` method is called in the subclasses constructors, so we could have chosen to make it a private method, but we chose to keep it public in case a user wants to change the `input_shape`, `neurons` and

other parameters, then they could call the method, and change, to fit other needs. In case the format of the pictures are different for example.

__repr__: The method `__repr__()` is also abstract here, because we do not want our objects that we use to be `NeuralNetwork`, we want them to be an instance of one of the two subclasses. Each subclass overwrites it to print a representation about what class they are.

_classifier: The `_classifier` is a private method, because it is only to be called in the `hidden_layers` method of the subclasses. when the `_hidden_layers` method is called.

The `_classifier` method creates an instance of the `Sequential` class based on the `_neurons`, `y_dim` and returns the `Sequential` object that is our classifier.

_call: The `_call` method is a private method to be used inside the `train` method. We do not need to call it outside the class. Therefore we chose to make it private. In our `_call` method we use the instance variable `_hidden`, `hidden` is different depending on what kind of instance it is. But the `_call` method can calculate the loss function for both instances.

train: The `train` method is just the same implementation as in the assignment text, we did not need to make any choices of how to implement it here.

test: The `test` method also uses the `_hidden` instance variable, so works for both `FullyConNN` and `ConvNN`. Here we also just followed the recipe in the text and did not have to make any assumptions.

Own methods: We have also added methods `set_hidden()`, `set_cls`, `set_params` to show how you can access and change the values of the private instance variables. We use these methods in the subclasses to change the initial values from `None` to the respective values we want, depending on the Subclass.

The `FullyConNN` class:

Constructor: In our constructor method here, the first thing we do is to change the value of our instance variable `_hidden` by calling the `hidden_layers` method. Then the `_hidden` is referring to a `Sequential` object. Then we set the value of the instance variable `_params`. This is to make the object fit to this class's specific needs.

__repr__: The method `__repr__` here is implemented to print out the class name for this class. So overrides the Super class implementation. Usually a `__repr__` returns a string representation, but the assignment specifically asked us to print out the string. Therefore we used a `print`.

hidden_layer: The method `hidden_layers` is implemented to fit the `FullyConNN` class. with `input_shape = 28*28` as default, and `neurons = 50` as default. We then make the `hidden` object and change the instance variable `_neurons` to the same value as the `neurons` parameter in the method, in this way we make sure that the `_neurons` is equal to the `neurons` in the `hidden_layers`. In the end we actually call the method `_classifier`. When we do this, every time we call the `hidden_layers` method, the `_classifier` method also gets called so that we don't have any mismatch between the number of neurons in `_hidden` and `_cls`.

The class inherits the instance variables from the `NeuralNetwork` class, The methods `__repr__` and `hidden_layers` are overridden by this subclass to meet the needs of this class. While it inherits the same implementation of the methods `train`, `test`, `set_hidden`, `set_cls` and `set_params`.

The `ConvNN` class:

Constructor: In this constructor method we also call the `hidden_layers()` method to set the values of the instance variables. Then we set the value of the instance variable `_params`.

__repr__: The method `__repr__` here is implemented to print out the class name for this class. Also here we used a `print()`, even though it is natural to return a string.

The `_hidden_layers` method is implemented differently here from the other subclass, made to fit this class needs. Here the default is `input_shape = (32, 32, 3)`, `neurons = 50`, `filters = 32`, `kernel_size = 3`, `strides = (2, 2)`.

The class inherits the instance variables from the `NeuralNetwork` class, The methods `__repr__` and `hidden_layers` are overridden by this subclass to meets the needs of this class. While it inherits the same implementation of the methods `train`, `test`, `set_hidden`, `set_cls` and `set_params`.

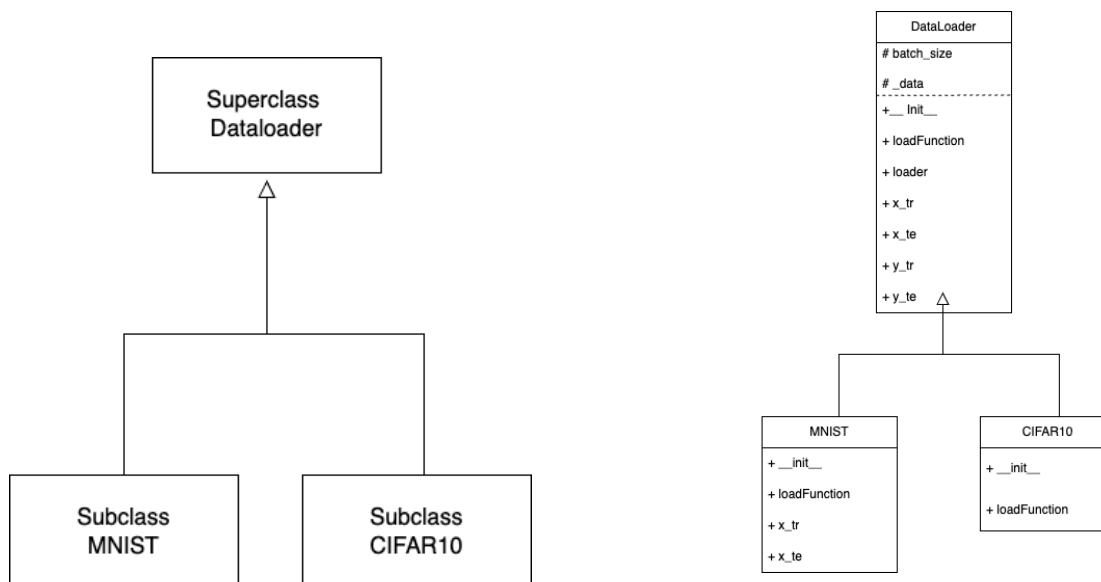
In the designed hierarchy, `NeuralNetwork` is a subclass of `tf.keras.Model`. `FullyConNN` and `ConvNN` are subclasses of `NeuralNetwork`.

While `tf.keras.Model` serves as the base class of our implementation we refer to `NeuralNetwork` as the superclass, because it is the direct parent class to `FullyConNN` and `ConvNN`. Overriding happens when `FullyConNN` and `ConvNN` provide their own implementation of a method inherited from the parent class. In this case we see it in the `hidden_layers` method. Even though `NeuralNetwork` might declare this method, the subclasses provide their own implementation by setting up layers according to their needs. Polymorphism makes it possible for us to use the same interface (call, `train`, `test`) on the different types of neural networks we have. For example, during model training, `train` will process data according to the specific architecture of the `FullyConNN` or `ConvNN` instance.

Task 3 DataLoader Class

3.1) Inheritance Diagram and Public Interface:

Instance variables
+ Public methods/ properties



In the `DataLoader` superclass we have methods like the constructor `__init__`, an abstract data loading method 'loadFunction' and a method 'loader'. The properties in this class are `x_tr`, `x_te`, `y_tr` and `y_te` which is training and test data.

In the MNIST subclass that inherits from DataLoader we have the constructor method and the LoadFunction, the properties are x_tr and x_te. We override the loadFunction method to be able to read the MNIST data set. It overrides the properties x_tr and x_te, to reshape the format. It inherits the loader method, and the properties y_tr and y_te.

In the CIFAR10 subclass the loadFunction method is implemented to read in the CIFAR10 dataset. while all of the properties are inherited from the superclass. The same for the loader method.

Both subclasses call the loadFunction in its constructor, in this way the _data instance variable changes value depending on what instance is created.

In summary, the public interface for these classes provides a consistent way to initialize the data loader, load the data, and access processed features and labels for both training and testing. The MNIST and CIFAR10 classes provide specific behaviors for these actions tailored to their respective datasets.

3.3) Inheritance and Overriding:

The superclass DataLoader is created to be generic enough to work with any dataset. It does not implement dataset-specific loading mechanisms, which is why “loadFunction” is abstract. MNIST and CIFAR10 subclasses inherit from DataLoader and gain access to its methods and properties, meaning they can use the loader method.

The MNIST and CIFAR10 classes override the loadFunction method of DataLoader. This means that both MNIST and CIFAR10 can provide their own implementations of loadFunction. In the MNIST class we also override the x_tr and x_te to reshape the data specifically for the MNIST dataset.

Task 4 Training Your Neural Network

We made a train/test-program where we assert that the input parameters are right. And then create correct objects depending on if the input is ConvNN or FullyConNN. And the same for CIFAR10 and MNIST. We then check that the model can calculate on the data. We assume that the different neural networks subclasses are specified to work on its own data set. And that is why we set default values that match the corresponding data set in our code.

If everything is okay, we just followed the recipe in the task for how to calculate the AUC score.

If the user wants to see the docstring for a specific class the user needs to write python3 train.py -doc <class name>.

To get help, user writes -h or -help.